

Build with Containers

Seoul, South Korea



Your presenters



- Held engineering and leadership roles at various silicon valley startups, enterprise organizations and consulting firms
- 3 years at AWS – Solutions Architect specializing in container adoption and container services business development

Omar Lari,
Container Service

Excited to be in South Korea!



- My first time here
- Suggestions for food, sights and soju drinking tips are welcome 😊

Today's Agenda

- 09:30 **Welcome**
- 09:35 Building with Kubernetes on AWS
- 10:20 Amazon EKS Hands-On Lab
- 12:45 **Lunch**
- 13:15 Partner Presentation
- 13:45 Building with Fargate and Amazon ECS
- 14:30 Fargate Hands-On Lab
- 16:30 **Closing Discussion**
- 17:00 End

Important information

- Bathrooms
- Slides will be provided after the event via slideshare
- You will receive a survey after the event in your email. It only takes **3 minutes** to complete! Let us know how we did and what to improve.

Containers on AWS

AWS Container Services Landscape

Application Management



Amazon Elastic Container Registry



Deployment, Scheduling, Scaling & Management of containerized applications



Amazon Elastic Container Service



Amazon Elastic Container Service for Kubernetes

Where the containers run

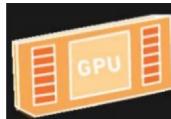


Amazon EC2



AWS Fargate

Where the containers run



Kubernetes on AWS

Open source container management

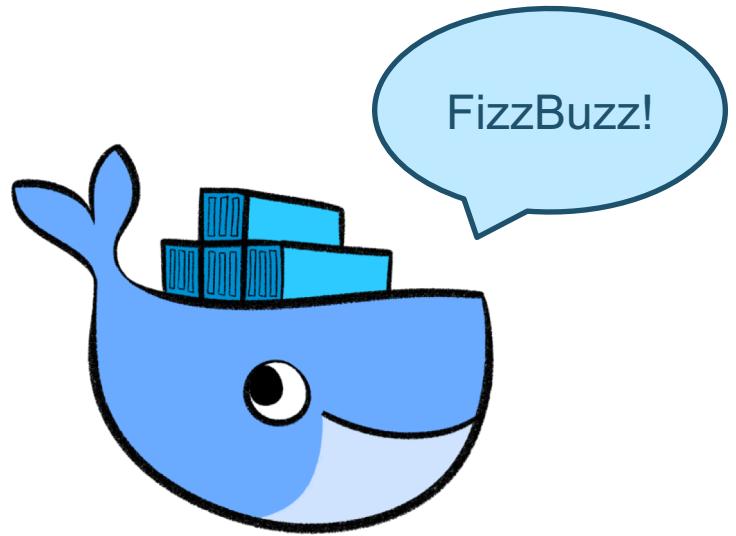
Omar Lari, Container Services

March 4th, 2019

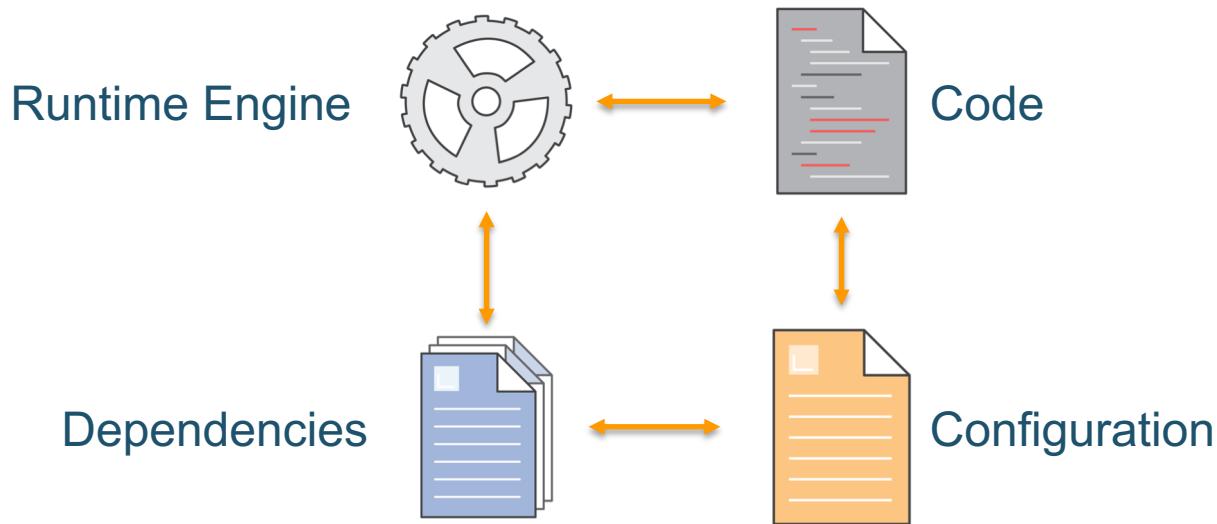
Agenda

- Why containers?
- What is Kubernetes?
- Key concepts
- Running Kubernetes on AWS
- Demonstration

Why containers?



Application environment components



Different environments



Local Laptop



Staging / QA



Production



On-Premises

It worked on my machine, why not in prod?



v6.0.0



v7.0.0



v4.0.0



v7.0.0



Local Laptop



Staging / QA

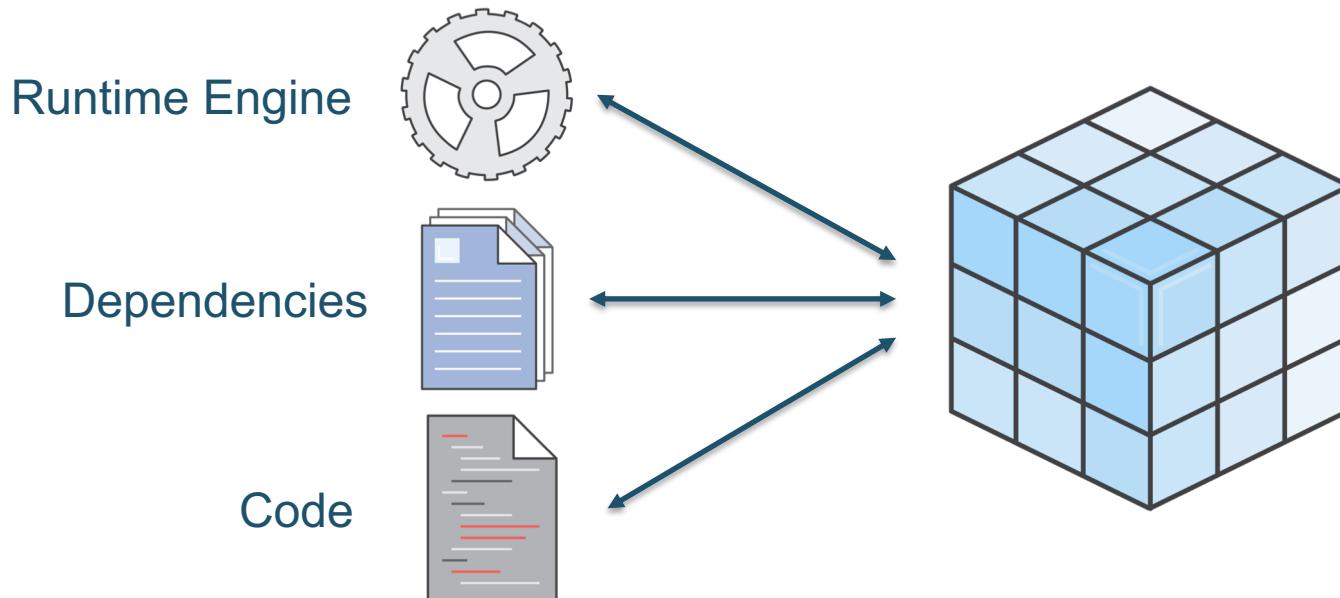


Production



On-Prem

Containers to the rescue



Docker

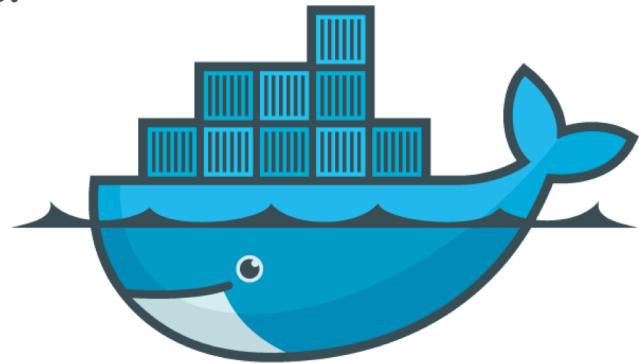
Lightweight container virtualization platform.

Tools to manage and deploy your applications.

Licensed under the Apache 2.0 license.

First released March 2013

Built by Docker, Inc.

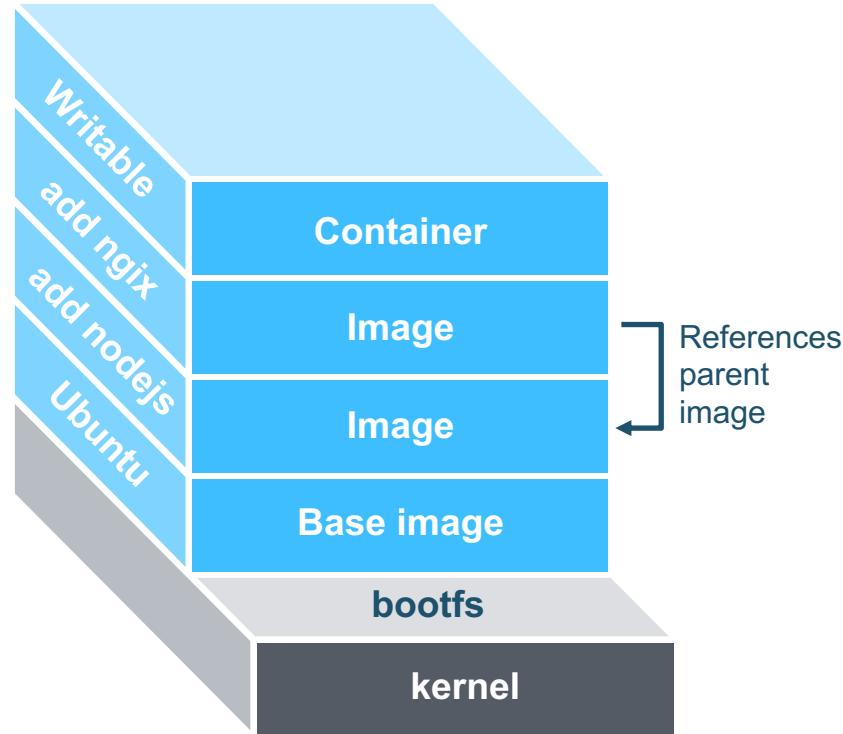


Docker Image

Read only image that is used as a template to launch a container.

Start from base images that have your dependencies, add your custom code.

Docker file for easy, reproducible builds.



Four environments, same container



Local Laptop



Staging / QA

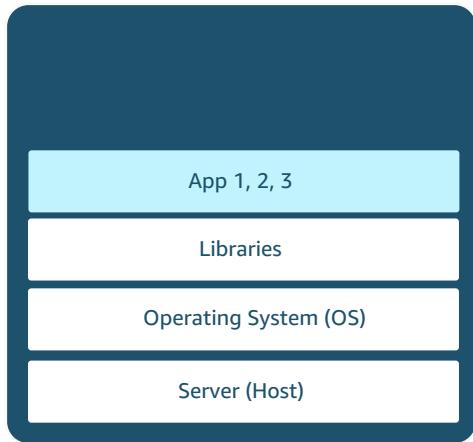


Production

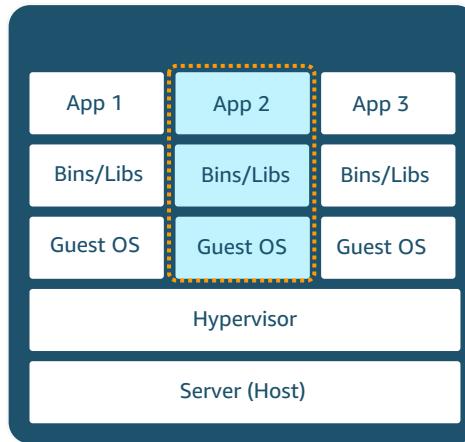


On-Prem

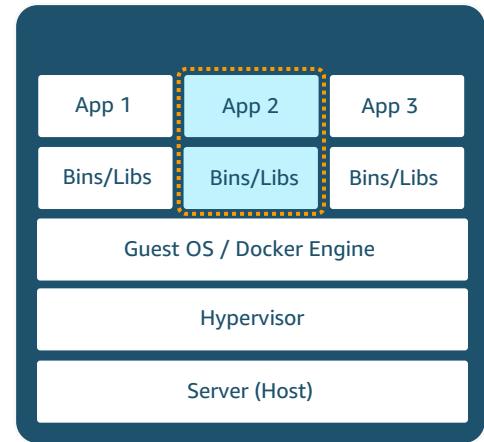
Container Implementation



Bare Metal



Virtual Machine



Containers

But how do we make this work at scale?



We need to

- start, stop, and monitor lots of containers running on lots of hosts
- decide when and where to start or stop containers
- control our hosts and monitor their status
- manage rollouts of new code (containers) to our hosts
- manage how traffic flows to containers and how requests are routed



kubernetes

aka 'k8s'



AmazonECS

AWS Container Services Landscape

IMAGE REGISTRY

Container Image Repository



MANAGEMENT

Deployment, Scheduling,
Scaling & Management



HOSTING

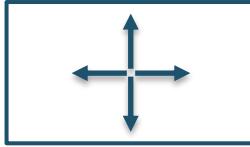
Where the containers run



What is Kubernetes?



**Open source container
management platform**



**Helps you run
containers at scale**



**Gives you primitives
for building
modern applications**

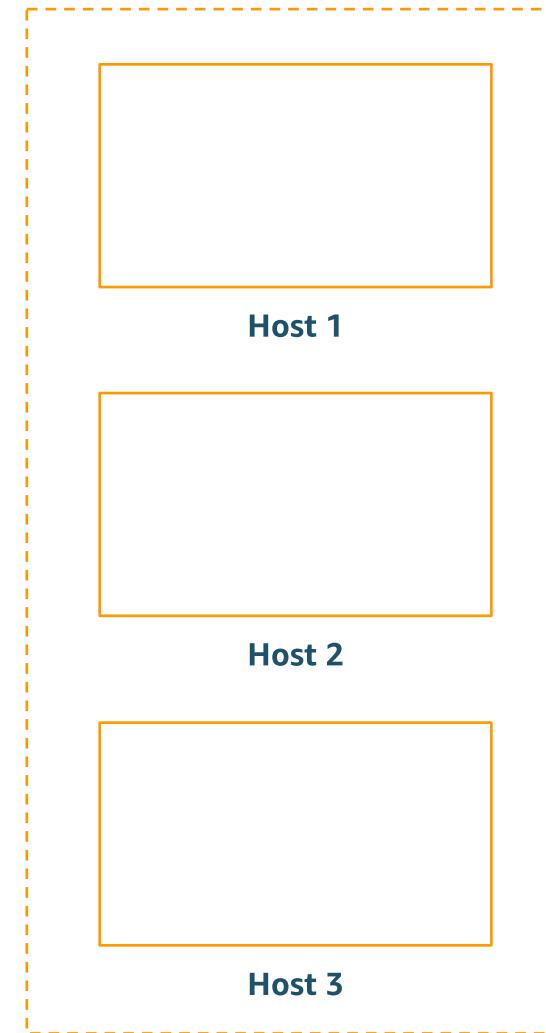
Containers on Hosts

A host is a server – e.g. EC2 virtual machine.
We run these hosts together as a cluster.

Our simple example web application is
already containerized.

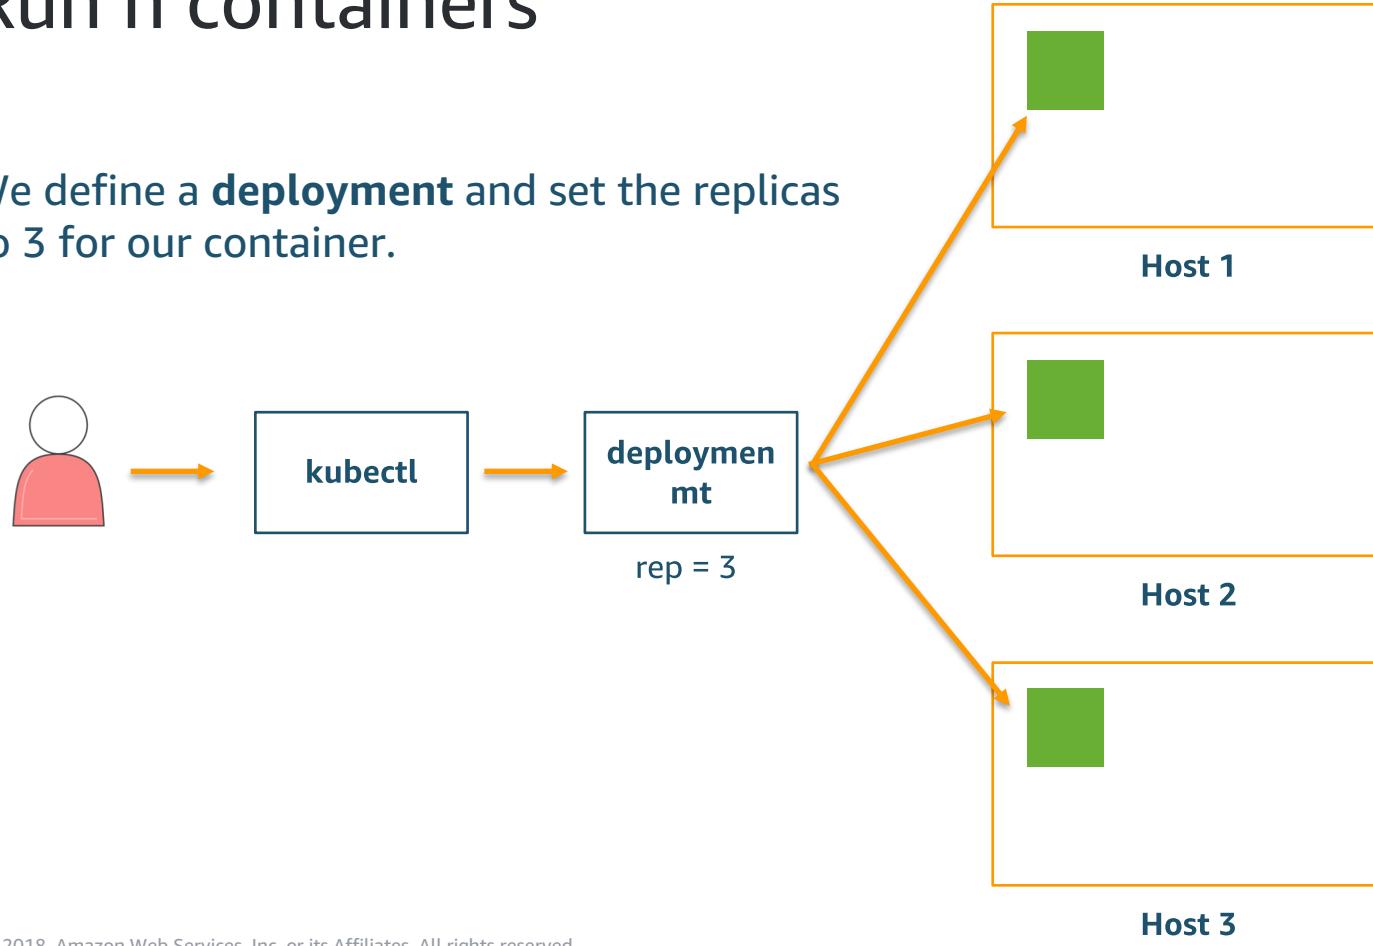


To start let's run 3 copies of our web
app across our cluster of EC2 hosts.



Run n containers

We define a **deployment** and set the replicas to 3 for our container.



Scale up!

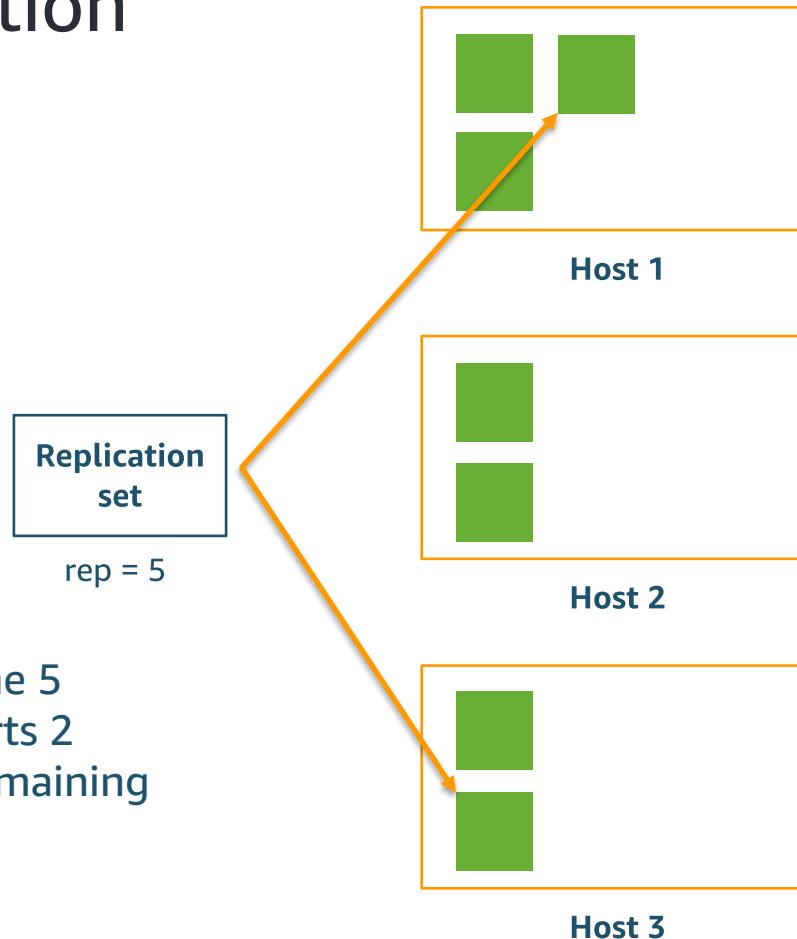
Need more containers?
Update the replication set!



The new containers are started on the cluster.

Untimely termination

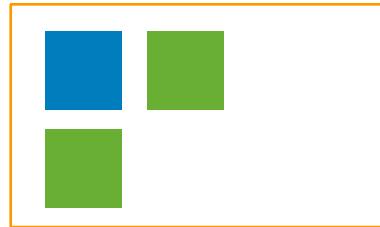
Oh no! Our host has died!



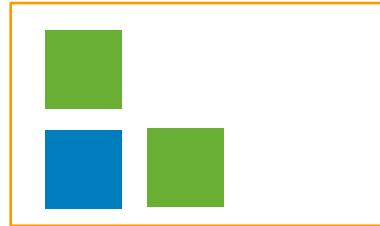
Containers IRL

In production, we want to do more complex things like,

- Run a **service** to route traffic to a set of running containers
- Manage the **deployment** of containers to our cluster
- Run multiple containers **together** and specify how they run



Host 1



Host 2



Host 3

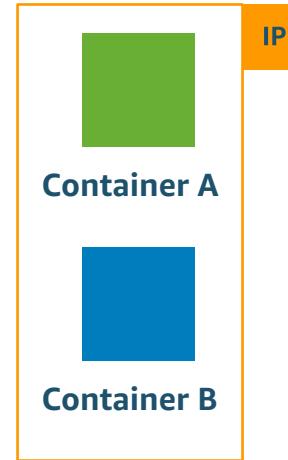
Pods

- Define how your containers should run
- Allow you to run 1 to n containers together

Containers in pods have

- Shared IP space
- Shared volumes
- Shared scaling (you scale pods not individual containers)

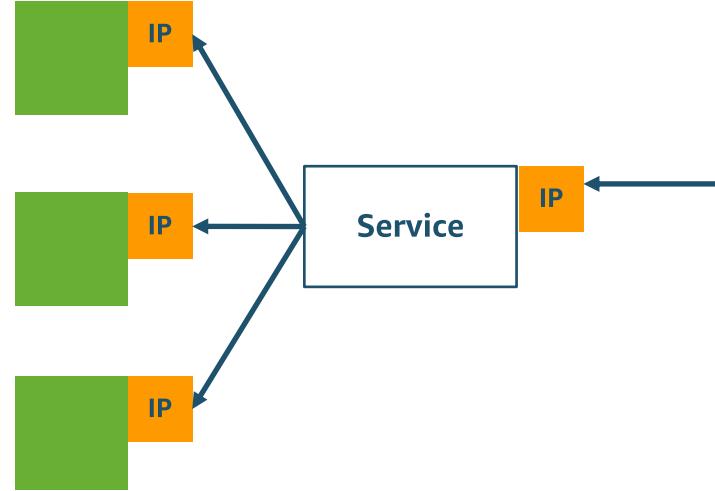
When containers are started on our cluster,
they are always part of a pod.
(even if it's a pod of 1)



Services

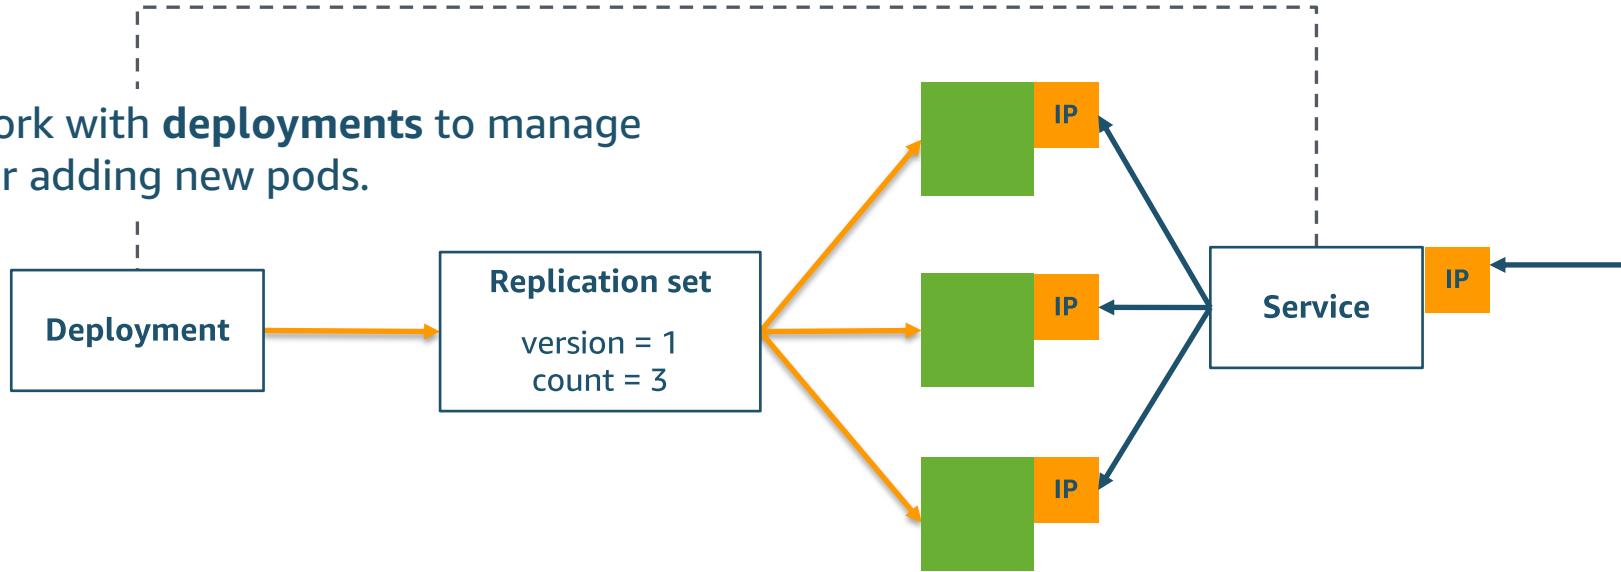
One of the ways traffic gets to your containers.

- Internal IP addresses are assigned to each container
- Services are connected to containers and use labels to reference which containers to route requests to



Deployments

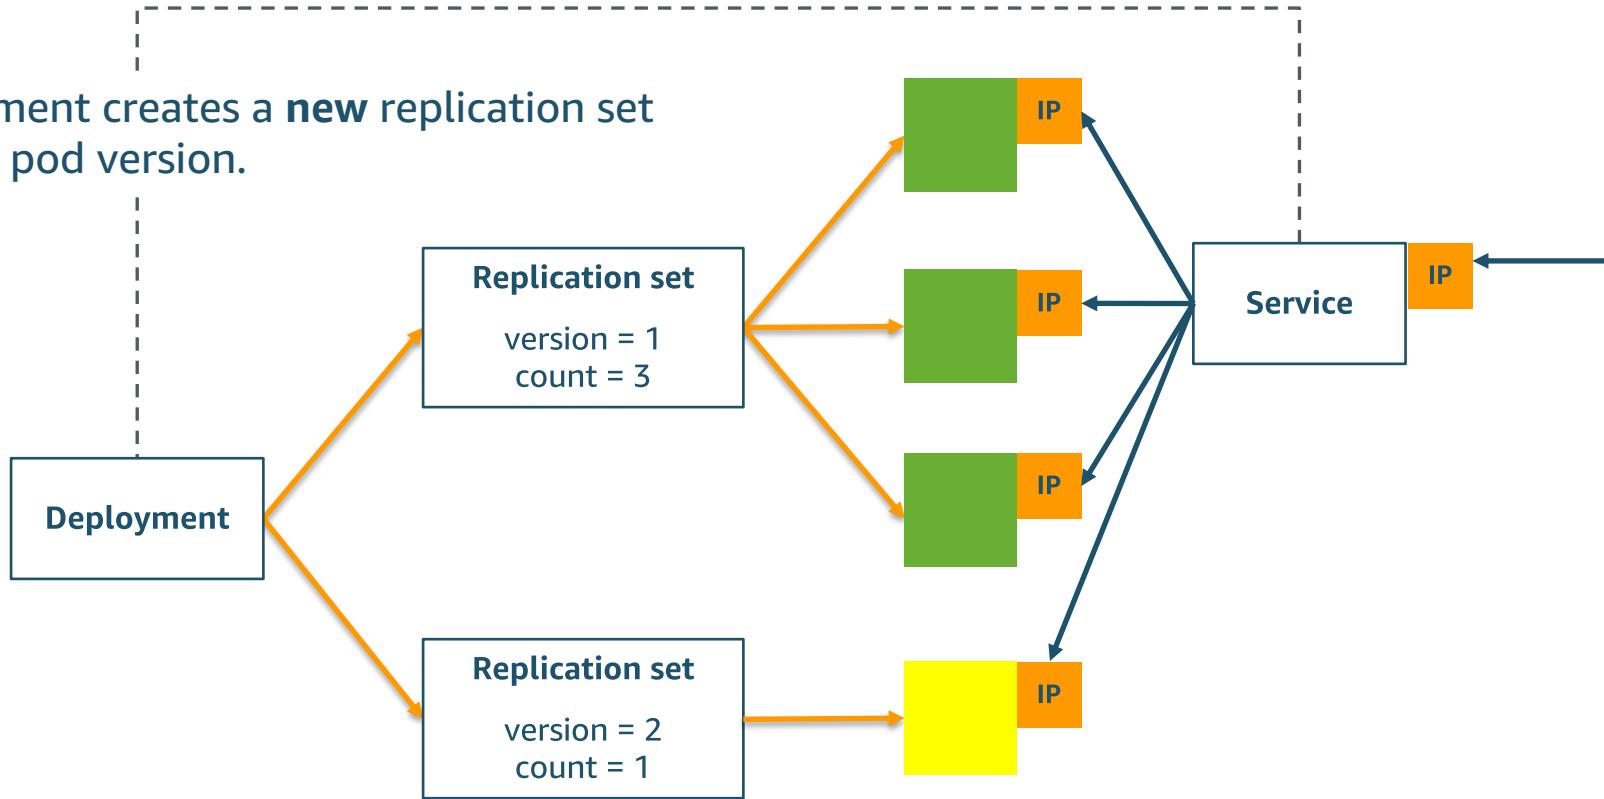
Services work with **deployments** to manage updating or adding new pods.



Let's say we want to deploy a new version of our web app as a 'canary' and see how it handles traffic.

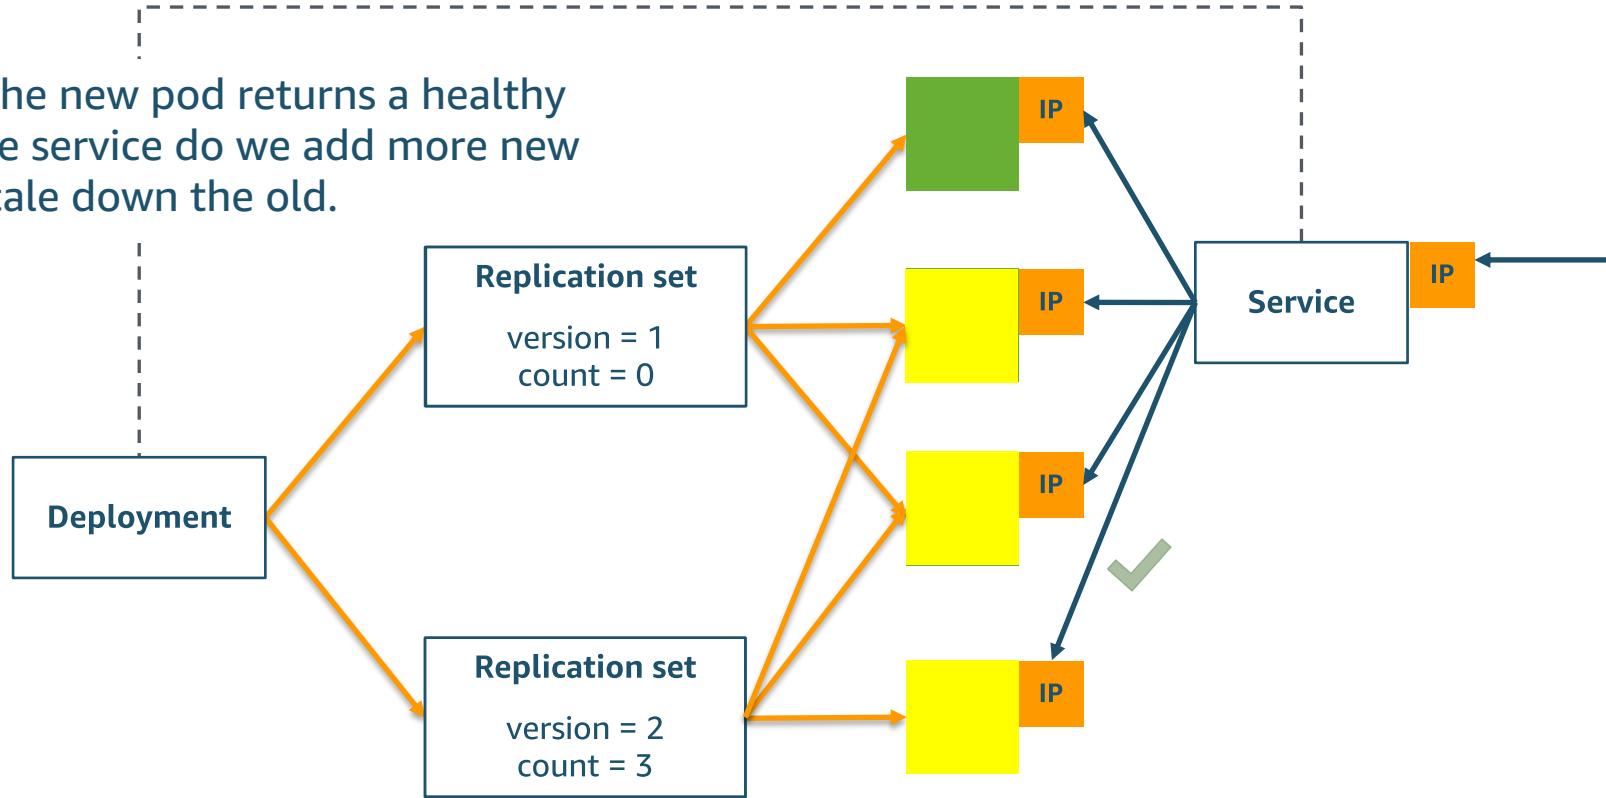
Deployments

The deployment creates a **new replication set** for our new pod version.



Deployments

Only after the new pod returns a healthy status to the service do we add more new pods and scale down the old.



Running Kubernetes on AWS





51%

of Kubernetes workloads
run on AWS today
— Cloud Native Computing Foundation

Elastic Container Service for Kubernetes



Managed K8s control plane — highly available API server and etcd nodes

Bring your own EC2 worker nodes

Core tenets

- Platform for enterprises to run production-grade workloads
 - Provides a native and upstream Kubernetes experience – Kubernetes certified
 - Seamless integration with AWS services
 - Actively contributes to the Kubernetes project
-

APIs

```
aws eks create-cluster \
--name <> \
--role-arn <> \
--resources-vpc-config <> \
...
```

eksctl – eks the easy way



CLI to provision and managed your EKS cluster

Quickly get started without worrying about VPC, IAM and Node design

eksctl create cluster

Or define a config file for more detailed configuration control

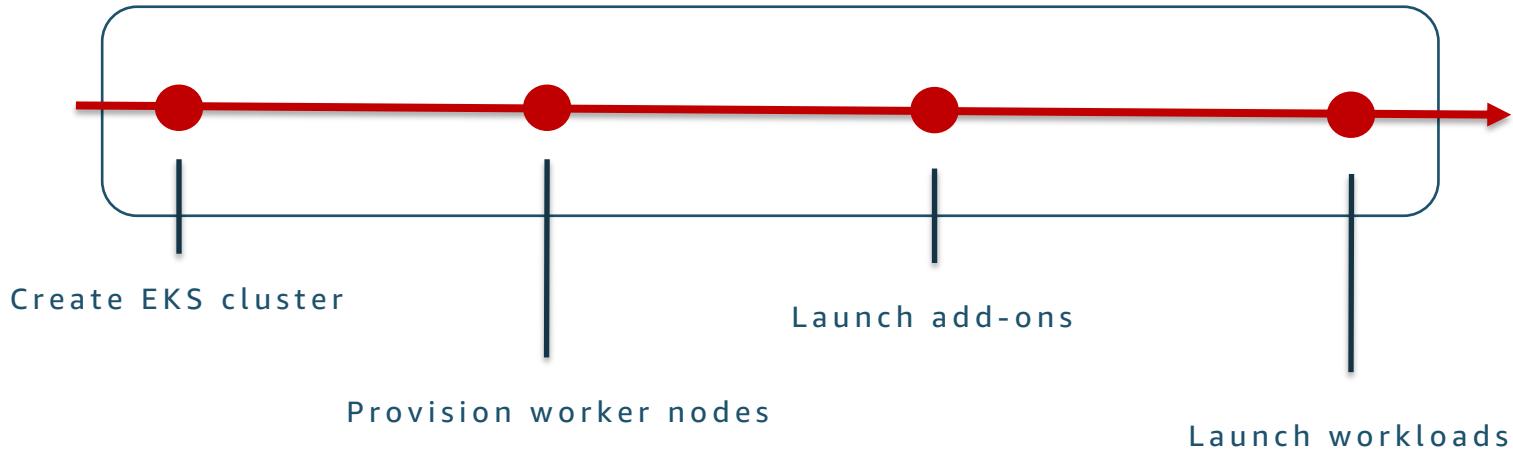
eksctl create cluster -f config.yaml

```
apiVersion: eksctl.io/v1alpha4
kind: ClusterConfig
```

```
metadata:
  name: demo
  region: ap-northeast-2
  version: '1.11'
nodeGroups:
  ...
```

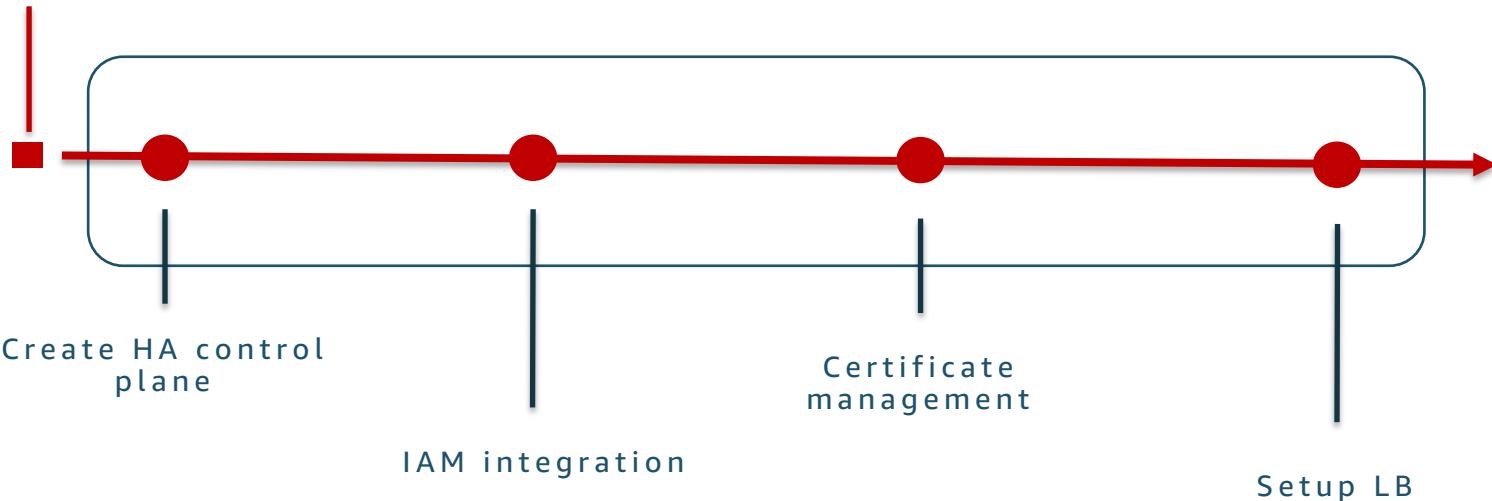


EKS Customers

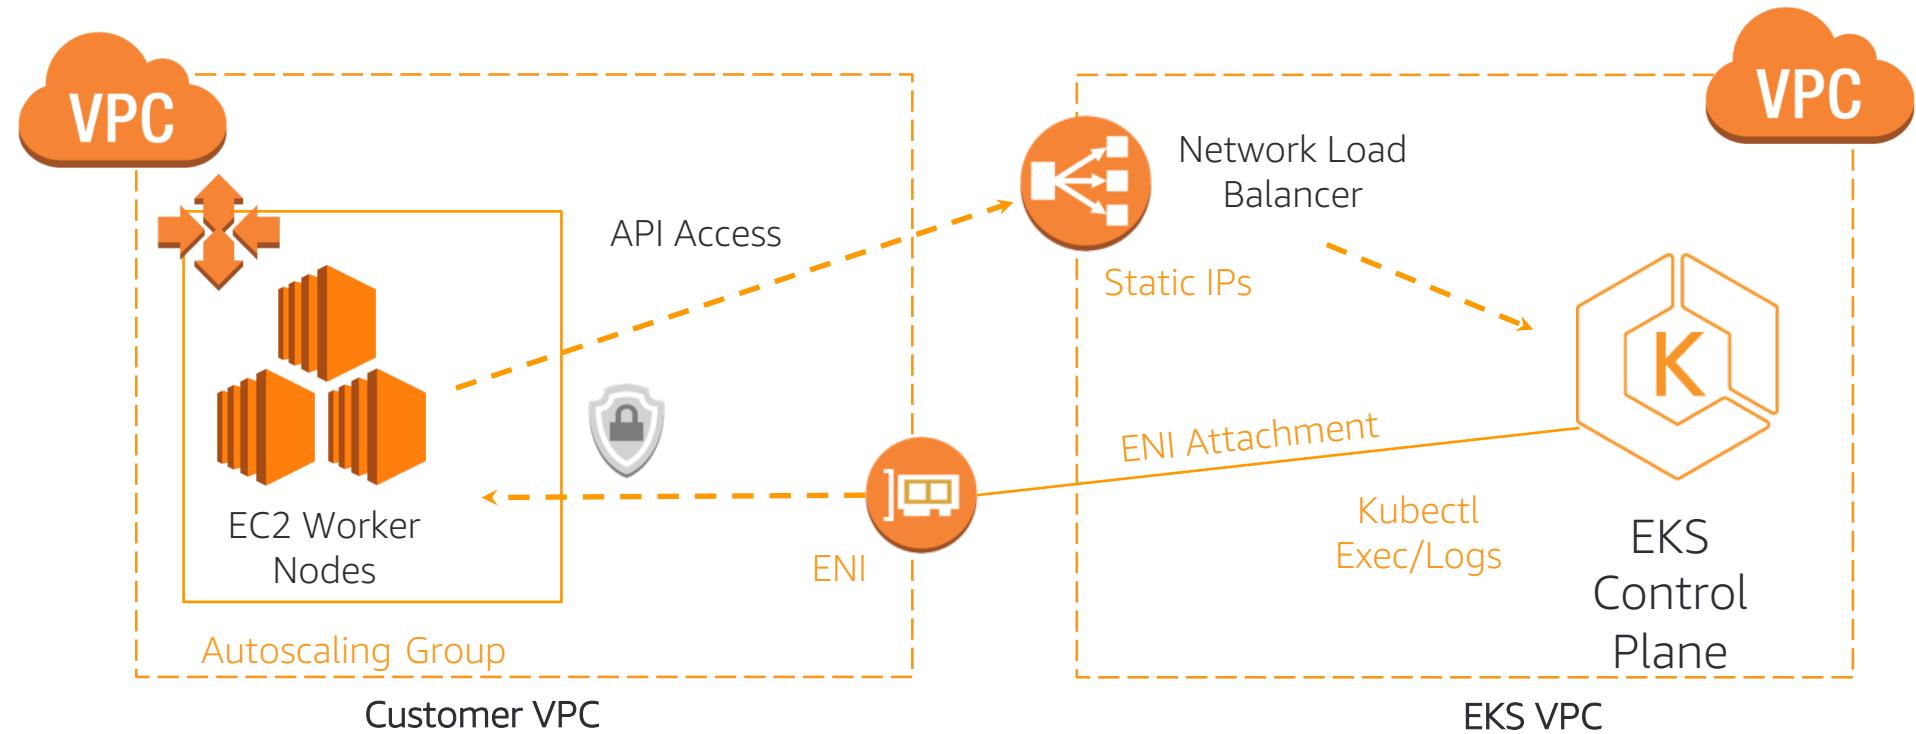


EKS – Kubernetes control plane

Create cluster



EKS Architecture





Container Networking Interface (CNI)

<https://github.com/aws/amazon-vpc-cni-k8s>



Native VPC networking
with CNI plugin



Pods have the same VPC
address inside the pod
as on the VPC



Simple, secure
networking



Open source and
on GitHub

VPC CNI plugin

- Bridge between the K8s land – AWS VPC
 - AWS Routable IPs
- Thin layer – no performance impact
- Pod IP ← ENI secondary IP

Cluster Identity & Access Management (IAM)



IAM enables secure access to AWS services and resources

Two IAM roles created
for K8s cluster

Service ▾	Access level	Resource
Allow (7 of 109 services) Show remaining 102		
Auto Scaling	Limited: List, Write	All resources
EC2	Full access	All resources
EC2 Container Registry	Full: List Limited: Read	All resources
ELB	Full access	All resources
ELB v2	Full access	All resources
Route 53	Limited: List	All resources
S3	Limited: List, Read, Write, Permissions management	Multiple

Service ▾	Access level	Resource
Allow (4 of 109 services) Show remaining 105		
EC2	Full: List Limited: Read, Write	All resources
EC2 Container Registry	Full: List Limited: Read	All resources
Route 53	Limited: List	All resources
S3	Limited: List, Read, Write, Permissions management	Multiple

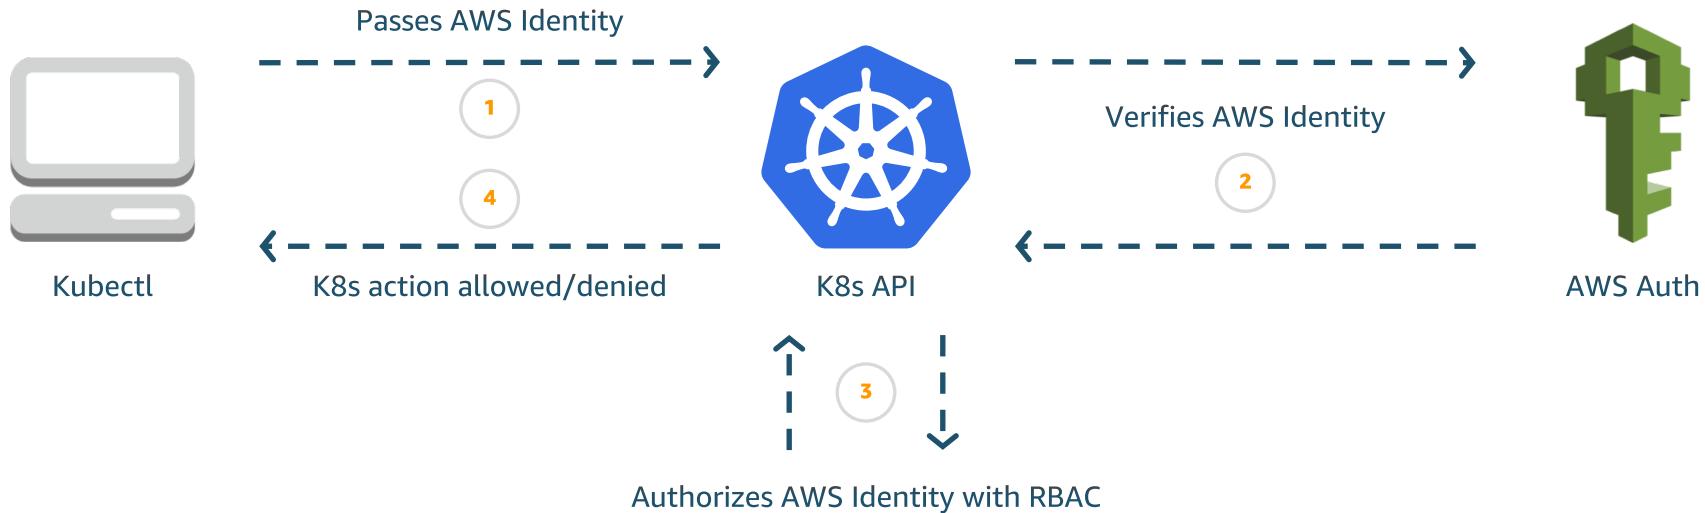
AWS IAM Authenticator for Kubernetes

- AWS native access management
- Built in collaboration with Heptio
- Kubectl and worker nodes
- Works with Kubernetes Role-Based Authentication (RBAC)

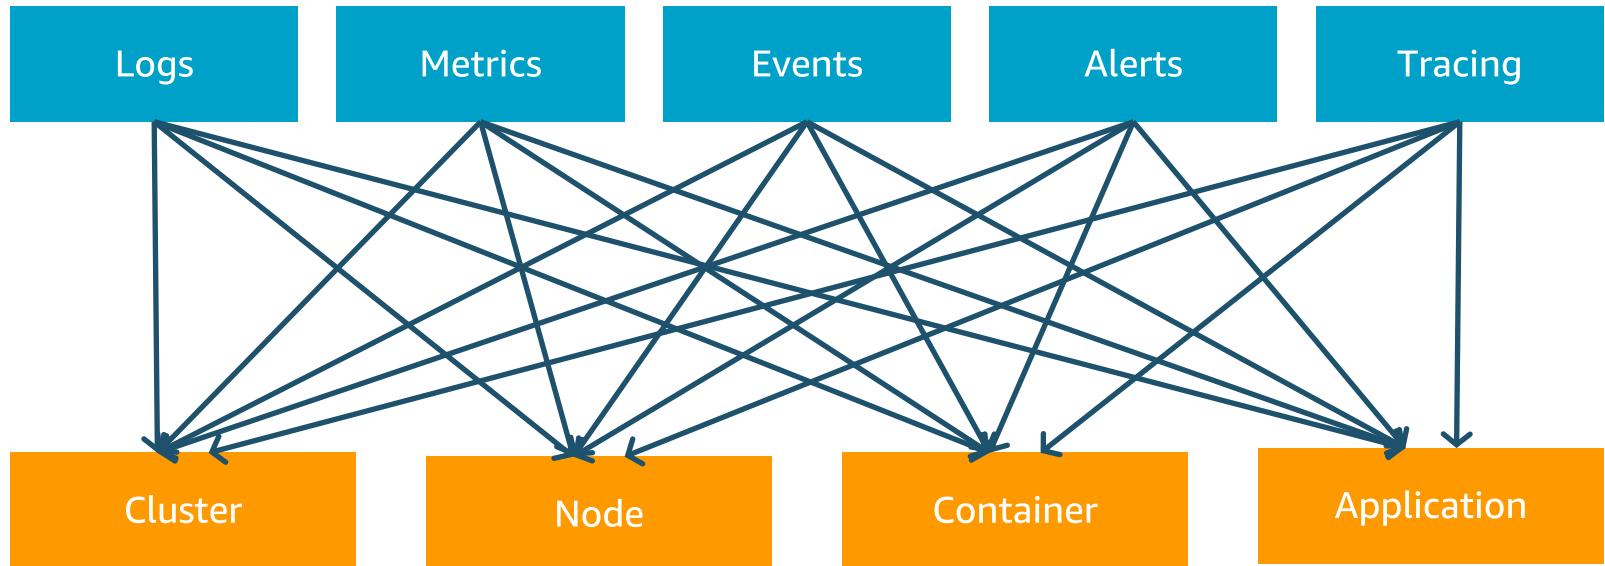


IAM roles for Kubectl

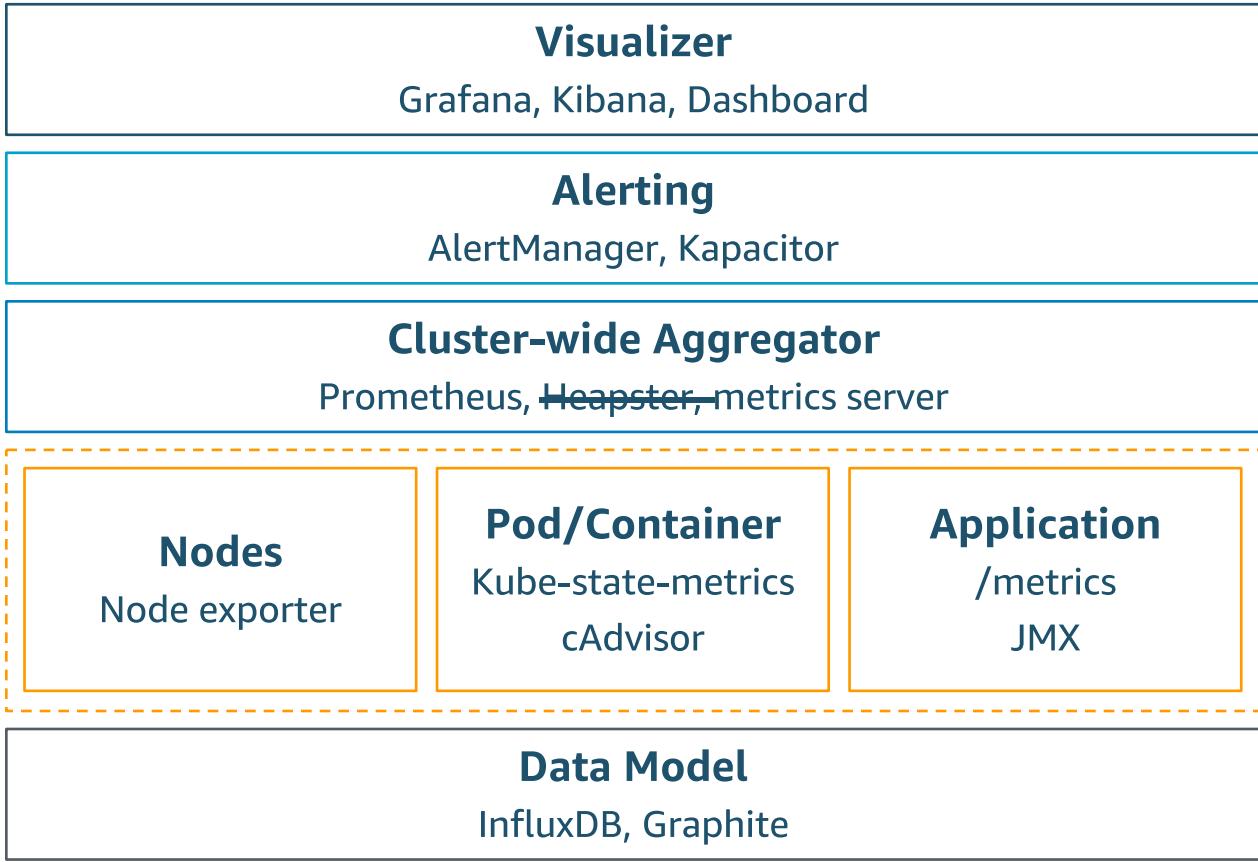
github.com/kubernetes-sigs/aws-iam-authenticator



Visibility throughout a Kubernetes cluster



Metrics



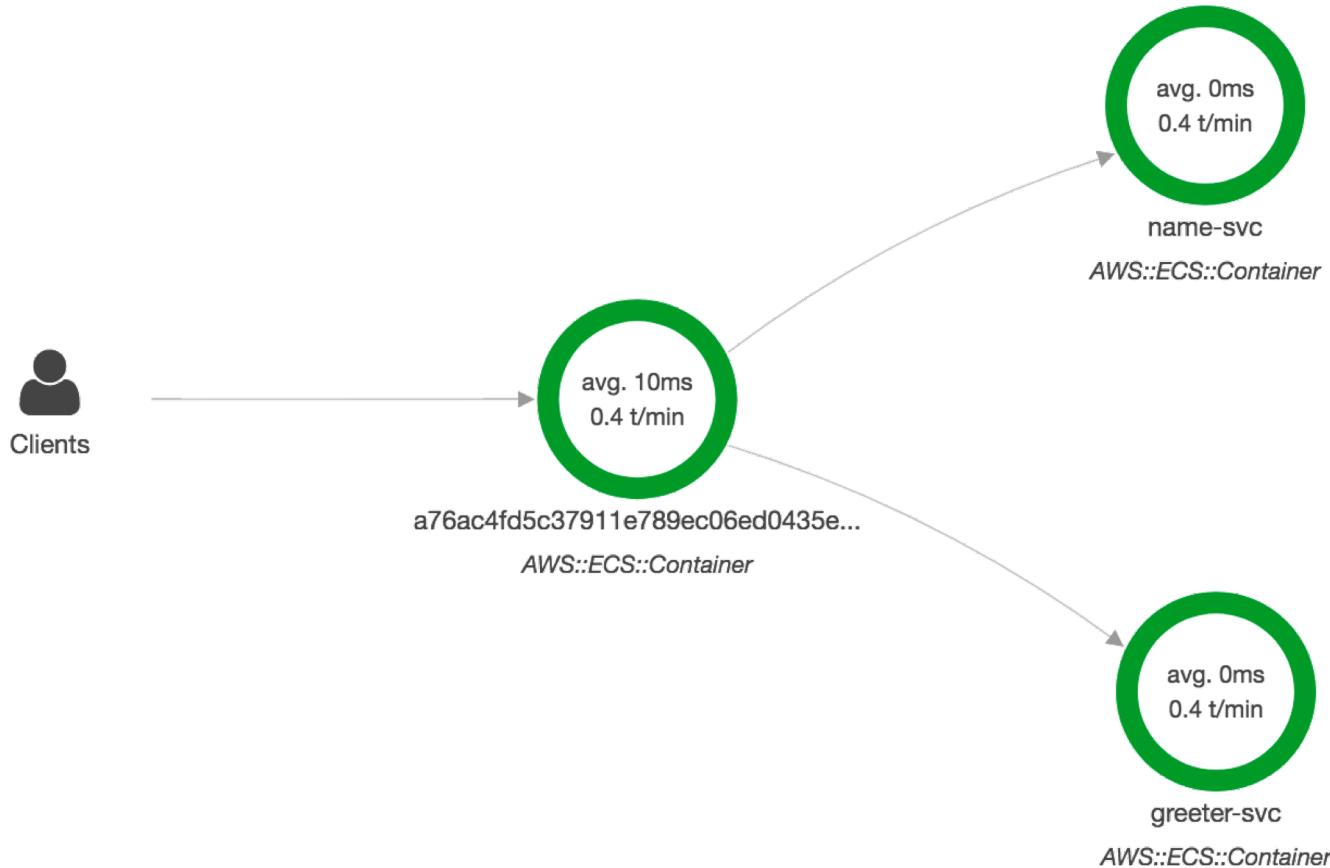
Tracing



Analyze and debug production, distributed applications, such as those built using a microservices architecture

Identify and troubleshoot the root cause of performance issues and errors

End-to-end view of requests as they travel through your application



CICD for Kubernetes Apps

CI/CD of apps on Kubernetes—choices

Jenkins

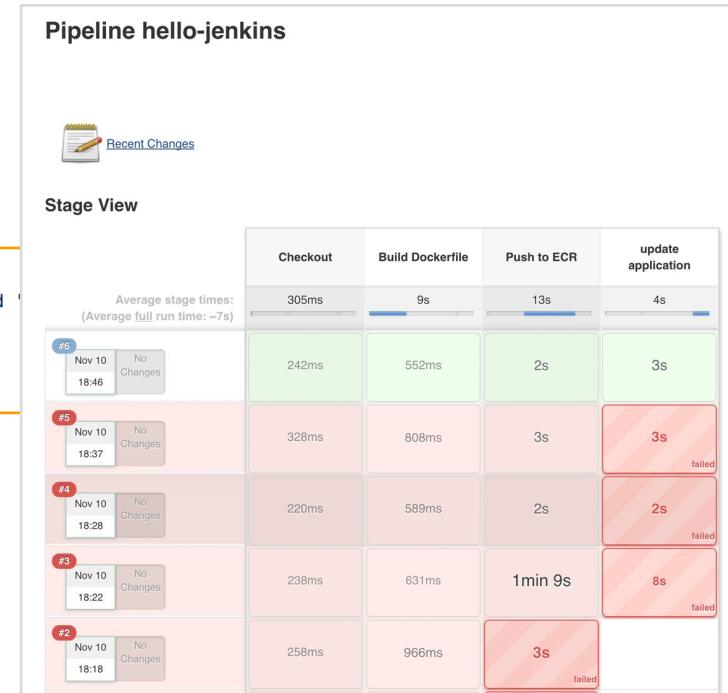
AWS CodePipeline, AWS CodeCommit, AWS CodeBuild

AWS partners

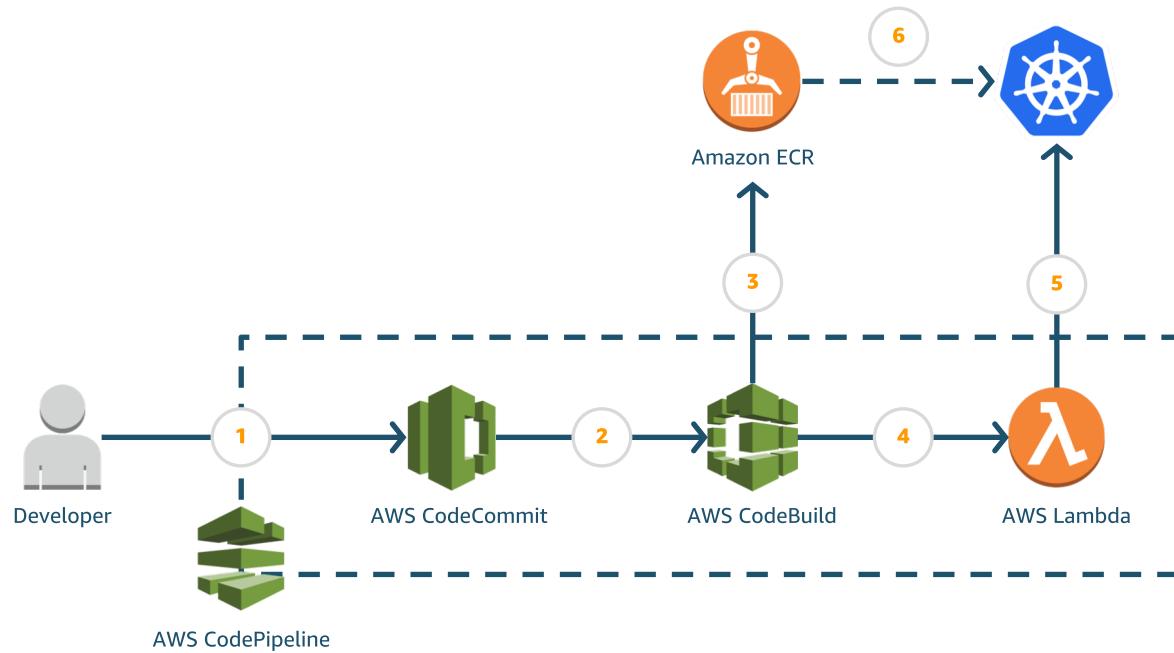
- GitLab
- Shippable
- CircleCI
- Codeship

Jenkins

```
1 node {  
2     stage 'Checkout'  
3         git 'https://github.com/omarlari/aws-container-sample-app.git'  
4     stage 'Build Dockerfile'  
5         docker.build('hello')  
6     stage 'Push to ECR'  
7         sh ("eval \$({docker run awscli aws ecr get-login --region ${REGION} --no-include-email | sed 's/^.*\. /$1 /')}")  
8         docker.withRegistry('https://${ECR_REPO}') {  
9             docker.image('hello').push('${BUILD_NUMBER}')  
10        }  
11    stage 'update application'  
12        kubernetes: { node {  
13            docker.image('kubectl').inside("--volume=/home/ec2-user/.kube:/config/.kube"){  
14                sh 'kubectl describe deployment ${APP}'  
15                sh 'kubectl set image deployment/${APP} hello=${ECR_REPO}/hello:${BUILD_NUMBER}'  
16                sh 'kubectl describe deployment ${APP}'  
17            }  
18        }}  
19    }  
20}
```



EKS Deployment Pipeline



- 1 Developers continuously integrate changes into a main branch hosted within a repo
- 2 Triggers an execution of the pipeline when a new version is found, builds a new image with build id
- 3 Pushes the newly built image tagged with build id to ECR repo
- 4 Invokes a Lambda function to trigger application deployment
- 5 Leverages Kubernetes Golang SDK to update a deployment
- 6 Fetches new container image and performs a rolling update of deployment

Questions?

Workshop

Start at <https://eksworkshop.com>

